Verified fixes. Honest acceptance status.

Release assessment | 21 October 2026 | Tool contract 14.0

The package has been repaired and tested locally. A 10/10 rating is not certified. The external evaluation, editorial review and deployed observation required by the supplied specifications are still missing. No numerical quality score has been invented.

Check	Measured result
Automated tests	332 passed; no failing tests. One third-party deprecation warning.
Corpus integrity	PASS: 43 indexed books, 6,935 pages, 95 qasidas and 2,682 canonical bayts.
Original source preservation	44 supplied raw source files checked against the uploaded ZIP; zero byte changes.
Disclosed PDF replay	28 calls, 0 errors; known hinted source page present in top five for 17/18 cases, compared with 7/18 in the uploaded baseline.
Local replay time	p50 377 ms; p95 1319 ms. These are local measurements, not Render latency.
HTTP / MCP boundary	Five tools discovered; 18 smoke checks passed. Eight simultaneous local calls: 0 errors.
Source and response contracts	Quotes slice exactly from canonical page text; compact and ultra projections pass evidence/provenance equivalence tests.

How to interpret the replay

A page-presence check is a narrow regression measure. It does not establish that the displayed passage fully answers the question, that the speaker attribution is editorially correct, or that unseen paraphrases work. The questions are disclosed development cases, not a blind test.

The uploaded baseline took about 15-16 seconds on two conditional questions. The revised planner recognizes those questions before the expensive generic search cascade. Current exact per-case timings and returned evidence are included in engineering_replay.json.

What changed

Mechanism changes with executable checks, rather than question-to-page routing.

Area	Implemented behavior
Retrieval	Persistent paragraph / Q&A; evidence units; source hashes and heading links. Bounded phrase, content, subject, lexical and concept routes are fused by route-family ranks. Candidate caps and route membership are visible in full detail.
Question intent	Frames retain the question type, content subjects, negation, qualifiers and alternatives. Longer signs, conditional practice, heart/work and Fawatih questions use source-local discovery vocabulary. Literal search/count paths retain their existing semantics.
Evidence and citations	get_page returns the canonical text used by search offsets. Typed hit validation checks source coordinates and exact quotes. Unreliable repeated printed labels are omitted; citations use the sequence.
Attribution	One explicit/structural speaker interpretation is applied across search modes. Full-passage roles and anchor roles are separate. Reported-opponent warnings survive both response projections.
Deadlines	A thread-safe checkpoint can be read by the supervising request. Later timeouts preserve validated lexical evidence; empty budget endings remain unknown. SQLite and stage boundaries cooperate with cancellation.
Diagnostics	Full detail includes stage events, route caps, candidate manifests, per-page filtering and named timing. Execution status is separate from evidence status. Retries preserve the original question and scope.
Operations	Build-tagged all-traffic attempt metrics are separate from scheduled probes. Dependencies are pinned. Startup rejects an old index missing the new evidence tables.
Client contract	Contract 14 exposes a reversible alias decoder and a shared typed source schema. Callable argument declarations drive advertised parameter lists.

Boundaries that remain explicit

The lexicon is a small, unreviewed development lexicon. This release does not claim a reviewed 600-concept morphology resource, a trained contextual reranker, calibrated support probabilities, or independently verified answer spans. Source-local scoring remains in use. These require domain judgments and class-level evidence before they can be accepted.

What your boss can rely on

Measured limits, deployment steps and ownership of final acceptance.

Token economy: measured, below the requested target

Control	cl100k: compact / ultra	o200k: compact / ultra
one hit question	809 / 762 tokens (5.81%)	735 / 688 tokens (6.39%)
six hit lexical	2421 / 2301 tokens (4.96%)	2063 / 1943 tokens (5.82%)
census	2273 / 2159 tokens (5.02%)	1924 / 1811 tokens (5.87%)
weak negative	656 / 647 tokens (1.37%)	641 / 633 tokens (1.25%)

Both encodings use tiktoken 0.12.0. Identical evidence and provenance are checked before comparison. Full workflow figures include tool discovery and a follow-up passage; no metadata hydration is required. The specified 25% reduction has **not** been demonstrated. Claude usage and unavailable client/system wrapper costs are not measured.

Remaining evidence for a defensible 10-tier claim

Independent evaluation: a frozen, privately owned question set with independently graded answer spans. The included stricter gate requires at least 1,002 wordings over 334 intent families, Direct@5 >=99%, Direct@1 >=95%, negative/weak controls and per-stratum checks. Public replay results cannot substitute for this.

Editorial and source review: Arabic/domain review of lexical links and support labels, at least 1,000 reviewed mixed hits, sensitive-speaker cases, and original scans for unresolved pagination and documentary provenance. No reviewer identity or review date was fabricated.

Production observation: deploy this build, refresh every MCP client, retain all-traffic metrics on durable storage, then collect at least seven days / 10,000 representative attempts. Local timing and ephemeral /tmp history cannot prove that gate. A Docker daemon and the user's Render account were unavailable, so no container build or production deployment is claimed.

Specific replay limitation: F13_LOVE did not return its hinted page in the top five. Other source candidates were returned; an independent reviewer must decide whether they are sufficient alternatives. No answer-grade credit has been assigned.

Deployment handoff

Use START_HERE.md. Rebuild the index, confirm /health reports 2.10.15 / contract 14.0, refresh clients, and run the supplied deployed smoke checker. Keep existing secrets in hosting environment variables. To finish acceptance, the team must supply the independent judgments, original-source review and production access/observation described above.

Boss-ready statement: "The connector has a tested engineering update with preserved sources and improved disclosed-case retrieval. It is ready for deployment verification, but a 10/10 quality claim remains unproven pending independent grading and production evidence."