Al Burhaniya Library 2.10.0 completion report

26 September 2026 | Server 2.10.0 | Tool contract 12.0

The updated source package fixes the main timeout path and the targeted ranking, citation, metadata and payload defects. All 218 regression tests pass. This report distinguishes implemented corrections from the external acceptance requirements in the supplied 2.9.0 developer report.

Measured local results

Check	Result
Regression suite	218 passed; 1 upstream deprecation warning
Report-question trials	240/240 completed; 100% in each tested family
Response-cache-cold latency	p95 0.66 s; maximum 1.96 s
Warm response-cache latency	p95 0.28 ms
30 bare phrases	p95 0.60 s; zero errors
HTTP concurrency	64/64 completed across 1, 2, 4 and 8 clients
Four-client / eight-client p95	3.93 s / 6.20 s
Lexeme census	14 literal pages within 20 lexeme matches; 0 missing
Two-book IDs payload	825 bytes, including timing and contract

Score status: 9+ not certified

No final numeric score is assigned. Local reliability thresholds now pass, but the PDF also requires independent retrieval grading, an unseen owner-held question set, client refresh tests, documentary evidence and seven days of production monitoring. Those requirements cannot be replaced by self-assigned scores. The release has not been deployed from this session.

Cold means the response cache was cleared; it does not mean the operating-system file cache was cleared. Five cold and five warm trials were run for each of 24 report questions. The first process call is included. HTTP results are local, on this execution host.

Reliability: every requested item

ID / status	Implementation and evidence
REL-1 Implemented; tested	Bounded stages retain earlier evidence with completion=partial. Forced later-stage timeout test passes.
REL-2 Implemented; tested	BM25 passage union capped at 300 windows; window scoring uses precomputed offsets and bounded text. Failing families complete locally.
REL-3 Implemented; tested	Bare phrases try exact and terms before bounded passage discovery; the slow legacy fallback is removed for short phrases.
REL-4 Implemented; measured	Named parsing, candidate, scoring, relation, excerpt and deduplication phases. Profile evidence included; microsecond overhead can exceed 10% on cache hits.
REL-5 Local acceptance passed	Real HTTP at 1, 2, 4 and 8 clients: 64/64 completed. Four-client p95 3.93 s; deployment hardware still needs replay.
REL-6 Implemented; tested	Bounded process cache keys include complete controls, corpus hash and file identity. Census/audit requests bypass response cache.
REL-7 Acceptance passed locally	Census retains 14 literal pages within 20 lexeme matches, zero missing. Timed evidence in acceptance.json.
REL-8 Implemented; operational gate open	Launcher starts five-minute monitored probes. Real seven-day production history has not elapsed; proxy/origin correlation still requires deployment review.
REL-9 Local schema passed; client gate open	Nine controls invoked through a fresh local MCP SDK session. Export and optional CLIENT_SCHEMA_DIGESTS startup guard included. ChatGPT/Claude refresh not performed.

Retrieval: every requested item

ID / status	Implementation and evidence
RET-1 Targeted acceptance passed	Subject phrases protected from trailing question verbs; Laylat al-Qadr answer ranks first in regression.
RET-2 Implemented; targeted checks	Co-occurring concepts and fixed-expression penalties suppress vanity/amazement and sermon-opener false matches. Broad homograph precision still needs independent grading.
RET-3 Acceptance passed	Attached and detached conjunction queries return the same three pages; relaxed spacing is declared.
RET-4 Implemented; broader gate open	Literal subject candidates are retained in the union; prefix-preservation regression passes. Best-phrase top-ten inclusion for every Appendix C question has not been independently certified.
RET-5 Acceptance passed	Near-duplicate snippets are grouped with also_in citations. W1/W2/W4 contain at most one biography copy.
RET-6 Targeted acceptance passed	W1/W2/W4 answer b41:131 ranks first. Local span and BM25 features break common ties; universal no-more-than-two-ties acceptance remains ungraded.
RET-7 Acceptance passed	Child-raising query returns only b44:29 and b44:35, marked weak; no root-drift filler.
RET-8 Implemented	Coverage register is published in the connector contract; coverage gaps are not invented away.
RET-9 Owner-dependent gate open	Runner accepts 40 independently owned questions over five forms. No unseen owner set was provided; no private accuracy is claimed.

Citation, provenance and metadata

ID / status	Implementation and evidence
CIT-1 Acceptance passed	Quoted opposing view at b07:256 carries reported_view and attribution_warning in both compact modes.
CIT-2 Implemented; scan review open	Errata register records the b39:5 typo and proposed reading; source text is unchanged. Parallel biography evidence flags it.
CIT-3 Preserved; scan review open	Sequence-only locators remain protected by regressions. b01 label 70 and b42 label 118 are recorded for original-scan review.
PROV-1 Partially resolved	Four source-supported metadata decisions documented: b05, b16, b17, b37. The broad automated inventory still flags 29 records; missing scans/colophons prevent the <=5 goal.
META-1 Implemented; grading gate open	Relevance is computed per hit and not inherited as a shared default. Targeted answer/biography tests pass. The required 100-hit human calibration is outstanding.
META-2 Implemented; grading gate open	Explicit Q&A markers and quotation cues replace broad dialogue/Quran guesses. Hadith-in-ornamental-brackets regressions pass. The 100-hit 95% precision gate is outstanding.
META-3 Implemented	ok means planned stages completed, degraded means partial processing, not_needed means literal path. fast/full/partial path is disclosed.
META-4 Acceptance passed	Article relaxation is declared article_variant; authoritative surface forms and safe/broad distinctions remain protected.

Diagnostics, payload and version clarity

ID / status	Implementation and evidence
DIAG-1 Implemented; tested	Tool errors include timing, stage, candidate counts and a concrete exact-query retry. Transport cancellation reports queue_or_worker_deadline.
DIAG-2 Implemented	Known issues retain documentary, independent grading, live deployment, client refresh and monitoring limitations. Local timeout measurements are reported separately.
TOK-1 Acceptance passed	Two-book IDs response is 826 bytes including contract and timing; repeated page-label inventory is only in full catalogue projection.
TOK-2 Implemented; tested	Shared pagination/default fields remain hoisted; individual relevance labels and citation warnings remain explicit.
TOK-3 Measured; exact E17/E18 unavailable	cl100k_base token counts recorded on supplied frozen identical-evidence fixtures. Literal formatting is unchanged from supplied 2.9 measurements; exact external E17/E18 response bytes were not supplied.
VER-1 Implemented	Release 2.10.0, contract 12.0; typo fixed, current instructions and full open-gate list included. Historical release evidence retained.

What the evidence does and does not establish

The public replay tests 30 familiar questions plus 11 report topics. Nonempty results and reference-page presence are not direct-answer accuracy. The included private-set runner requires 40 owner-held questions spanning definition, comparison, practice, colloquial and negation forms. It produces responses for independent grading; it does not invent grades.

The source review records exact excerpts, offsets, source hashes and before/after metadata for b05, b16, b17 and b37. Original Arabic source files remain byte-for-byte unchanged. The broad review inventory still flags 29 records, including unidentified compilers; this is broader than the PDF's 15-record authorship count. Neither count is represented as resolved.

Deployment and acceptance handoff

Deploy this package

1. Replace the backend project with the contents of Al_Burhaniya_Connector_2.10.0. Preserve your deployment secrets outside the archive.

2. Build with the included Dockerfile, or install requirements-dev.txt and run scripts/build_index.py, pytest and scripts/independent_corpus_audit.py. The generated database is deliberately excluded from the ZIP.

3. Start with python scripts/start_server.py. It honors PORT and defaults to two workers. Confirm /health reports 2.10.0 and contract 12.0; MCP remains /mcp/.

4. Refresh the connector tools separately in ChatGPT and Claude. Record the schema digest and run all nine controls in docs/CLIENT_ACCEPTANCE_2100.md. Local SDK success does not prove either client refreshed.

5. The launcher starts five-minute monitoring unless MONITOR_ENABLED=false. Use a persistent MONITOR_SUMMARY_PATH to retain seven days across restarts. Review failures and latency before signing off.

Required before a final 9+ assessment

Run the supplied deployment and load checks against the real endpoint, with first attempts retained.

Collect seven days of production monitoring with the required success rate.

Grade the owner-held private set and the required 100-hit relevance/type samples. Verify the broader phrase-preservation and tie conditions.

Obtain the missing original title pages, colophons and scans for authorship, pagination and errata decisions. Preserve unresolved attribution where evidence is insufficient.

Evidence included in the ZIP

verification/2.10.0 contains pytest.txt, acceptance.json, load.json, public_replay.json, payload_measurements.json, issue_matrix.json and the package manifest. data/provenance_decisions_2100.json and data/text_errata.json record source-review decisions. README.md is the current operational guide; older release reports are historical.